



# Using Packets to Guide Server Optimization

Josh Clark

Distributed Performance Engineer

[https://github.com/je-clark/sf24\\_eu\\_server\\_optimization](https://github.com/je-clark/sf24_eu_server_optimization)



- M.S. in Computer Engineering
  - Focus on networking, protocol design, and security
- Distributed Performance Engineer
  - Using packets to solve complex performance issues

<https://www.jeclark.net>

<https://github.com/je-clark>

[https://github.com/je-clark/sf24\\_eu\\_server\\_optimization](https://github.com/je-clark/sf24_eu_server_optimization)



## A brief history of the QUIC protocol and Google

<https://www.androidpolice.com/quic-protocol-guide/>

## An Argument for Increasing TCP's Initial Congestion Window

Nandita Dukkkipati, Tiziana Refice, Yuchung Cheng, Jerry Chu  
Tom Herbert, Amit Agarwal, Arvind Jain and Natalia Sutin  
Google Inc.

Mountain View, CA, USA  
{nanditad, tiziana, ycheng, hkchu, therbert, aagarwal, arvind, nsutin}@google.com

<https://research.google/pubs/an-argument-for-increasing-tcps-initial-congestion-window/>

## Snap: a Microkernel Approach to Host Networking

Michael Marty, Marc de Kruijf, Jacob Adriaens, Christopher Alfeld, Sean Bauer, Carlo Contavalli\*, Michael Dalton\*, Nandita Dukkkipati, William C. Evans, Steve Gribble, Nicholas Kidd, Roman Kononov, Gautam Kumar, Carl Maurer, Emily Musick, Lena Olson, Erik Rubow, Michael Ryan, Kevin Springborn, Paul Turner, Valas Valancius, Xi Wang, and Amin Vahdat  
Google, Inc.  
sosp2019-snap@google.com

<https://research.google/pubs/snap-a-microkernel-approach-to-host-networking/>

## How Netflix tunes Ubuntu on EC2



Jorge O. Castro

on 11 August 2015

Tags: [ec2](#), [Ubuntu](#)

<https://ubuntu.com/blog/how-netflix-tunes-ubuntu-on-ec2>



facebookincubator/katran

A high performance layer 4 load balancer

● C · ☆ 4.6k · Updated 1 hour ago

<https://github.com/facebookincubator/katran>

CLOUDFLARE The Cloudflare Blog

All Posts Product News Speed & Reliability Security Zero Trust Developers AI Policy Partners LI

## How to drop 10 million packets per second

07/06/2018



Marek Majkowski

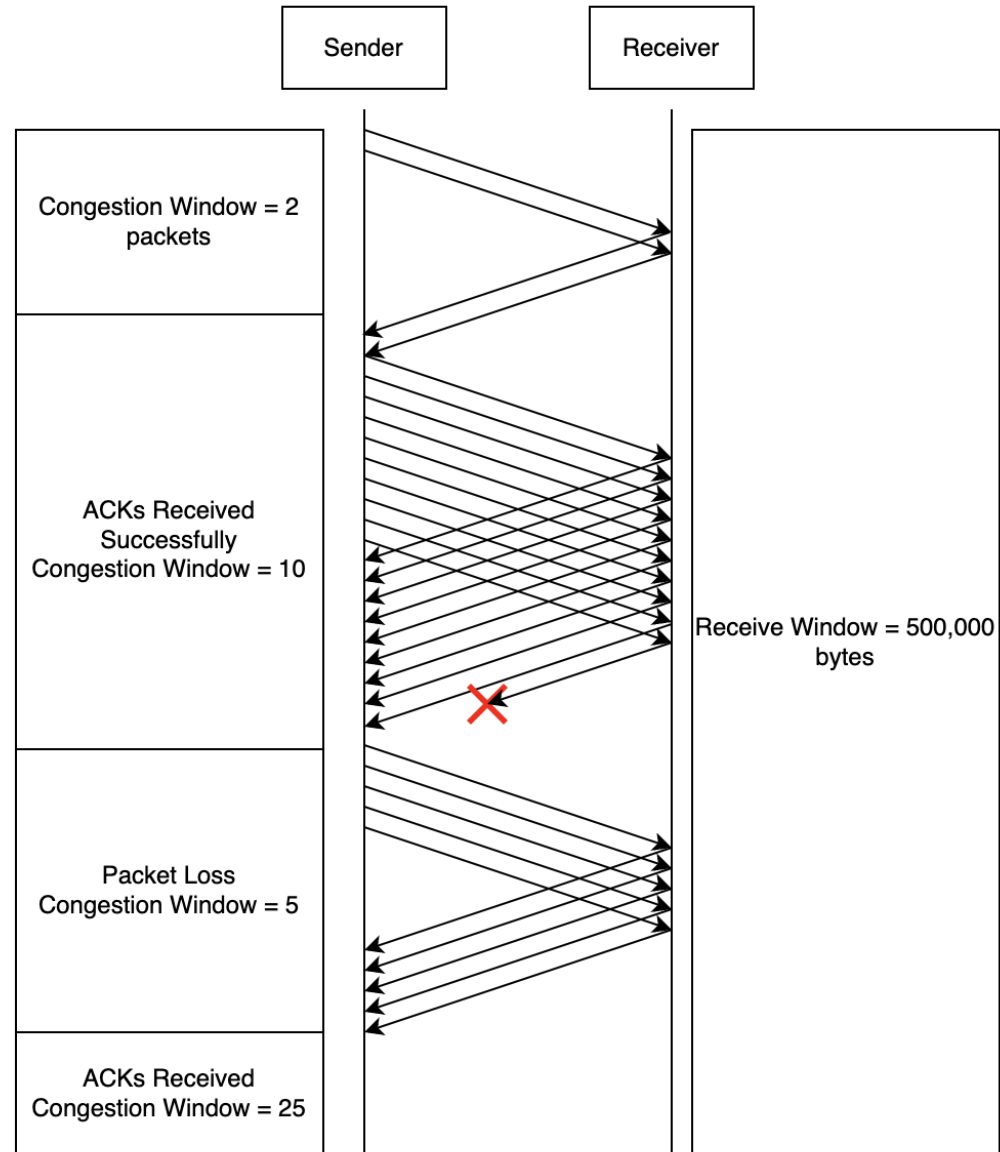
<https://blog.cloudflare.com/how-to-drop-10-million-packets/>



- TCP Windowing Review
- Baseline TCP Settings to Enable
- Evaluating and Optimizing Long TCP Flows
- Evaluating and Optimizing Short TCP Flows



# TCP Windowing Review





# Baseline TCP Settings




```
pi@server:~ $ sudo sysctl -a | grep tcp
net.ipv4.tcp_frto = 2
net.ipv4.tcp_sack = 1
net.ipv4.tcp_timestamps = 1
net.ipv4.tcp_window_scaling = 1
```





# Long TCP Flow Optimization



$$\text{BDP} = \text{bandwidth} \times \text{latency}$$



```
net.ipv4.tcp_rmem = 4096 131072 6291456
```

Minimum

Default

Maximum


- Test Network BDP = 625 KB
- Linux can use up to 3x more memory than the receive window value
- Receive memory should be at least 1.8 MB



0\_scp\_200MB.pcap



- Tracks a max window where loss occurs
- Logarithmic increase to max window, then cubic above max window
- Window  $\times 0.2$  when loss occurs




<https://www.cs.princeton.edu/courses/archive/fall16/cos561/papers/Cubic08.pdf>



- Existing algorithms use packet loss to measure congestion, but that isn't accurate
  - Packet loss can happen randomly without congestion
  - Congestion-induced packet loss is a lagging metric for congestion


Figure 1. Delivery rate and round-trip time vs. inflight.





- Estimates BDP using RTT variance
- Sends data at BDP rate
  - Periodically attempts a pacing\_gain to see if bottleneck bandwidth has changed
- Ignores small amounts of packet loss

Figure 4. RTT (blue), inflight (green), and delivery rate (red) detail.





1\_scp\_200MB\_bbr.pcap





2\_scp\_mystery.pcap



```
adding: files/file_477 (deflated 27%)  
adding: files/file_366 (deflated 27%)  
adding: files/file_963 (deflated 27%)
```

```
real    0m0.398s  
user    0m0.285s  
sys     0m0.098s
```

```
pi@client:~$ time scp -r pi@192.16  
pi@192.168.2.15's password:  
files.zip
```

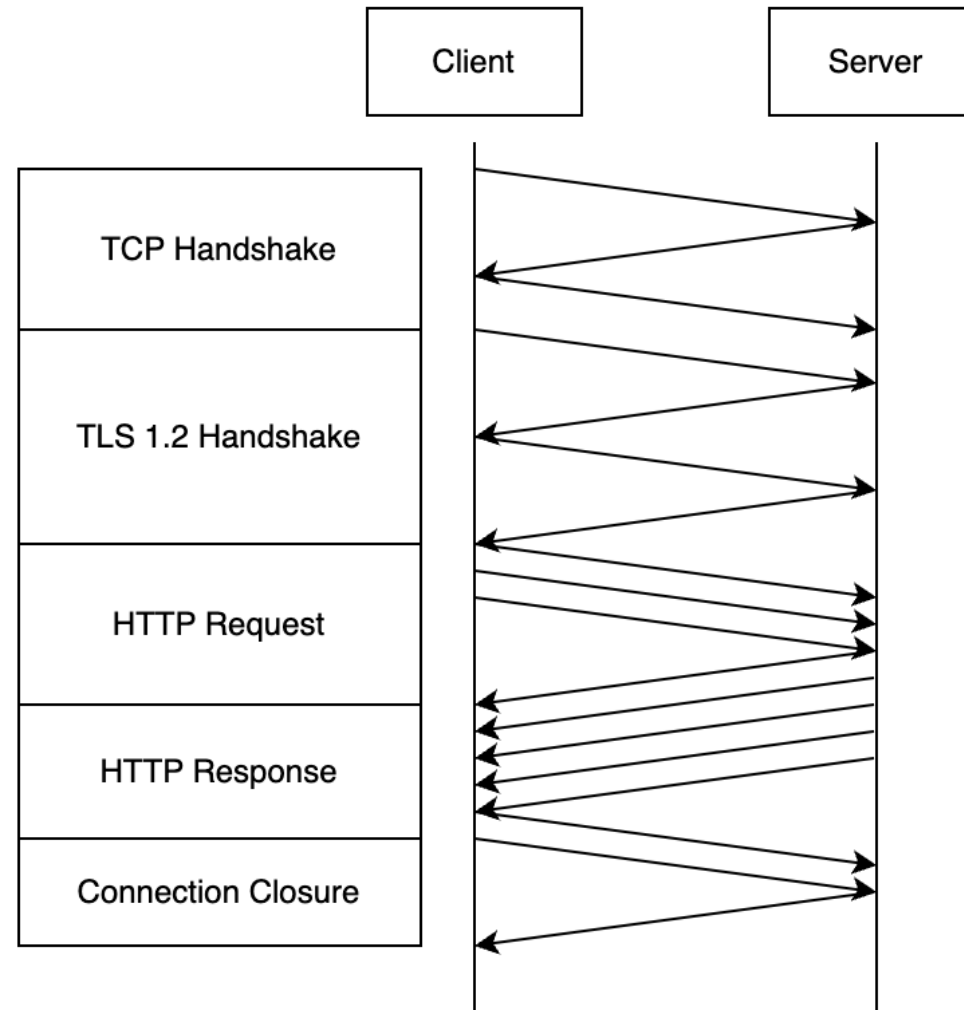
```
real    0m2.286s  
user    0m0.285s  
sys     0m0.069s
```

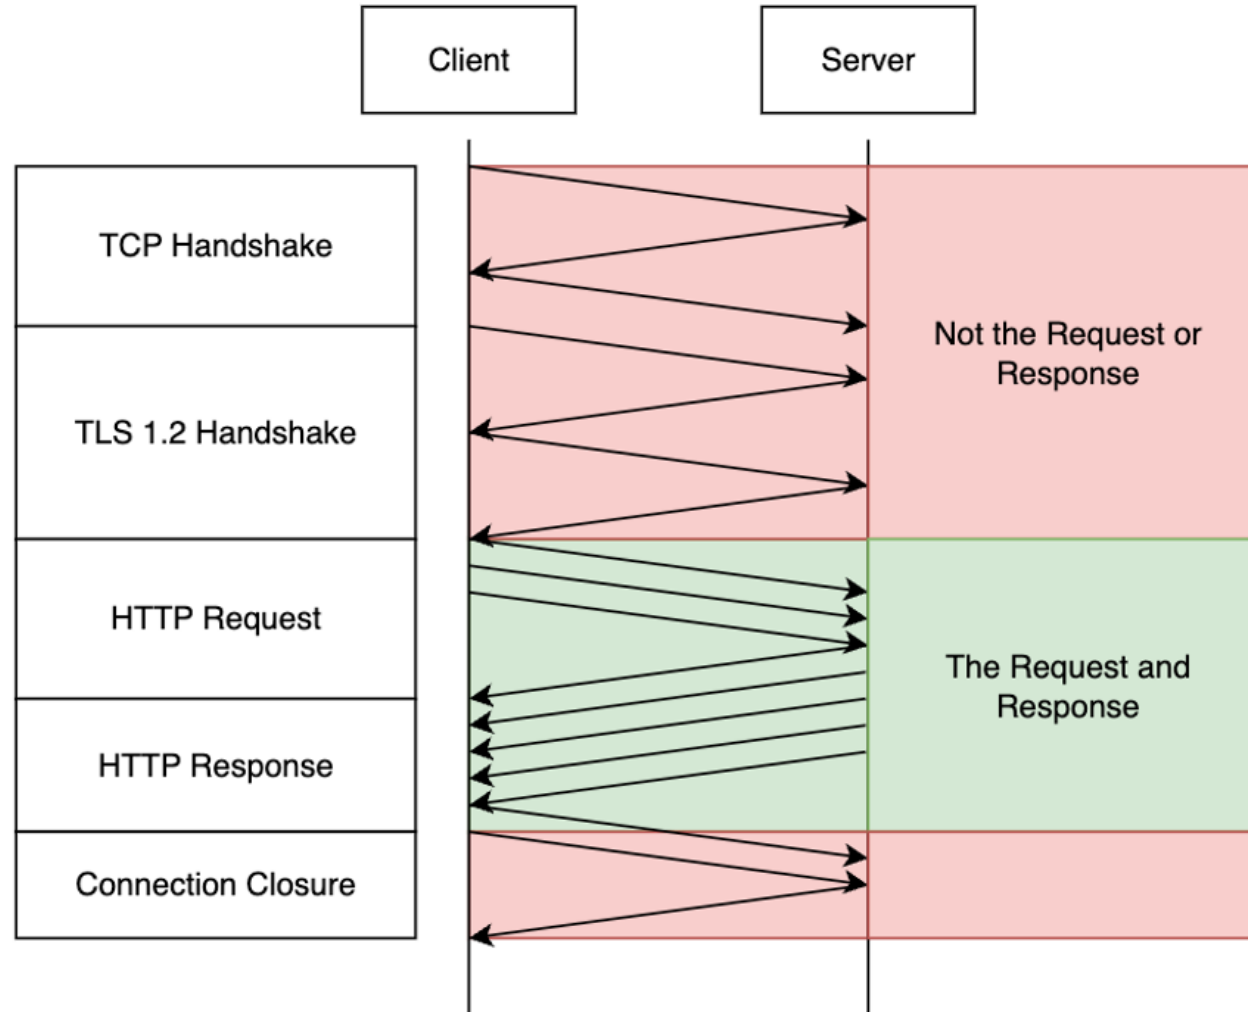
```
inflating: files/file_477  
inflating: files/file_366  
inflating: files/file_963
```

```
real    0m0.244s  
user    0m0.107s  
sys     0m0.122s
```



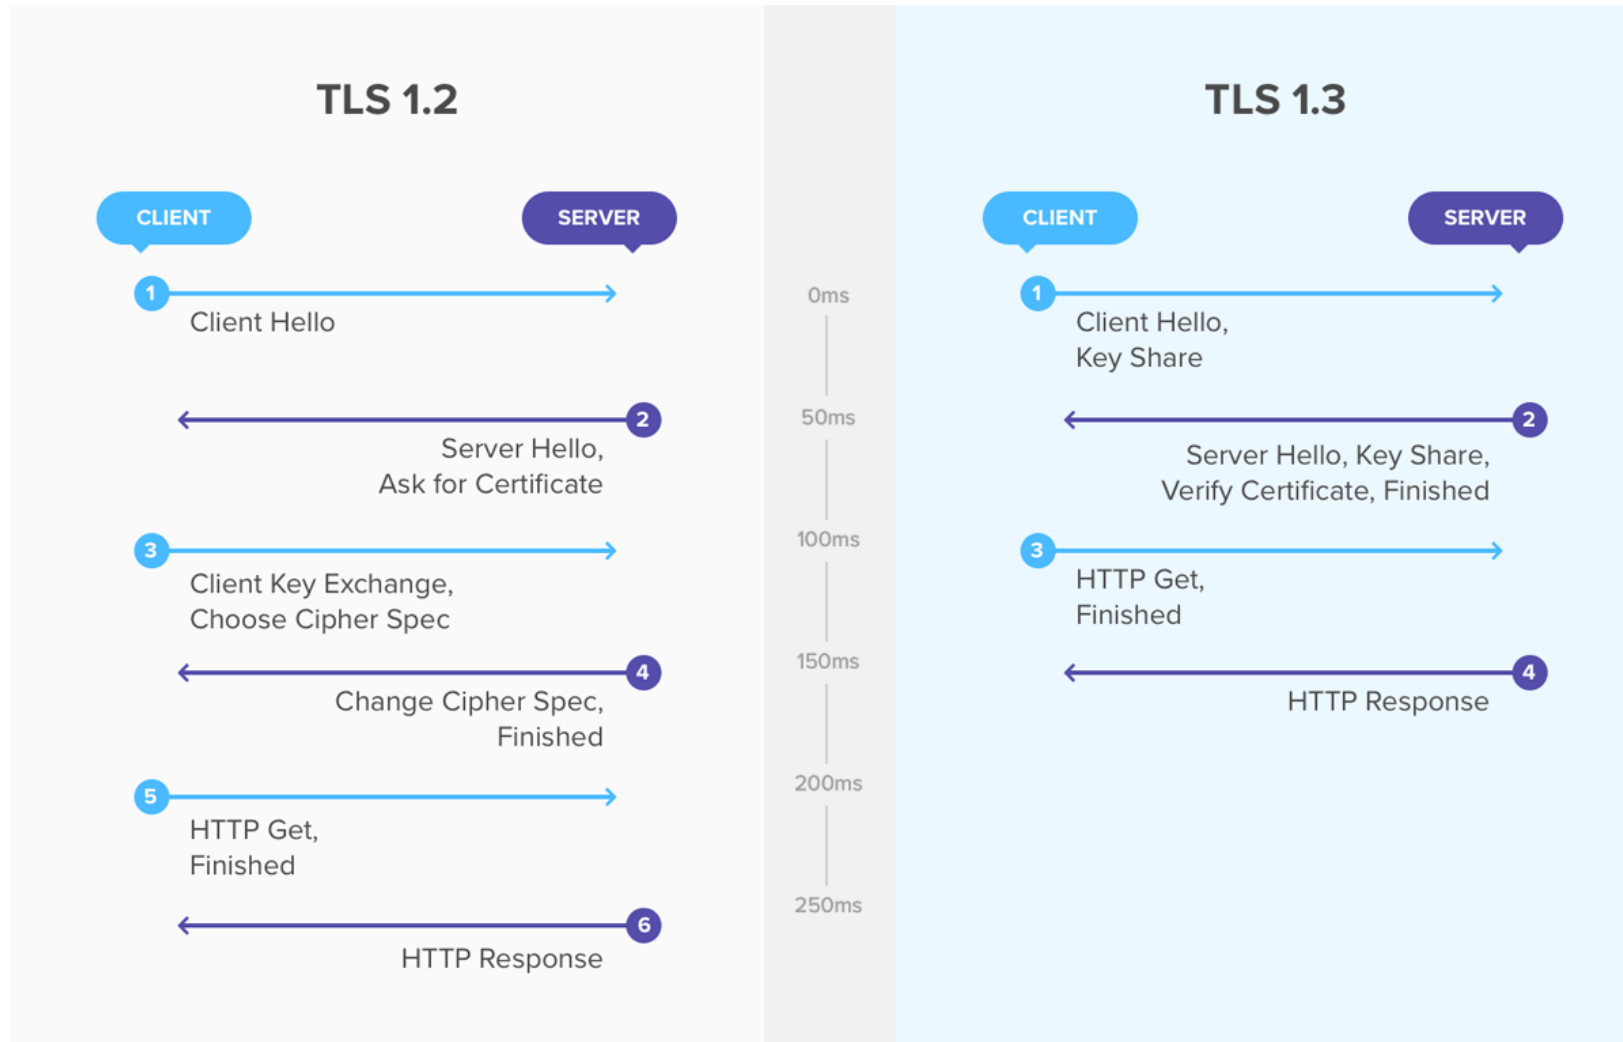
# Short TCP Flow Optimization







3\_small\_http\_request.pcap



<https://www.eyerys.com/articles/news/improved-internet-security-protocol-tls-13-has-been-approved>



4\_small\_http\_request\_tls13.pcap





- In 2011, Google changed `init_cwnd` to fit 90% of responses in 1 RTT

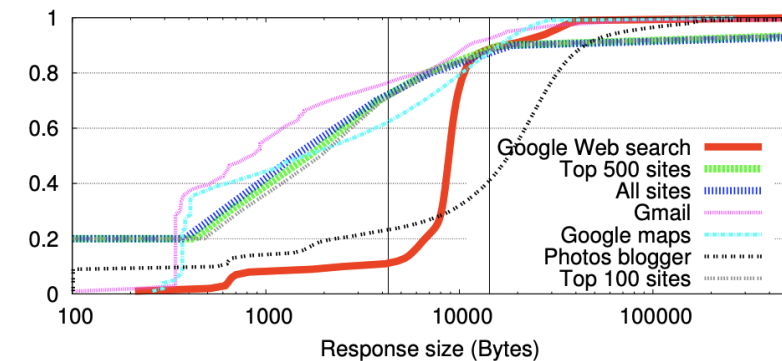


Figure 1: CDF of HTTP response sizes for top 100 sites, top 500 sites, all the Web, and for a few popular Google services. Vertical lines highlight response sizes of 3 and 10 segments.




Figure 2: TCP latency for Google search with different `init_cwnd` values.



```
pi@server:~ $ ip route show
default via 192.168.2.1 dev eth0 proto dhcp src 192.168.2.15 metric 100
192.168.2.0/24 dev eth0 proto kernel scope link src 192.168.2.15 metric 100
pi@server:~ $ sudo ip route change default via 192.168.2.1 dev eth0 proto dhcp src 192
.168.2.15 metric 100 initcwnd 25
pi@server:~ $ ip route show
default via 192.168.2.1 dev eth0 proto dhcp src 192.168.2.15 metric 100 initcwnd 25
192.168.2.0/24 dev eth0 proto kernel scope link src 192.168.2.15 metric 100
```



5\_small\_http\_request\_initcwnd\_45.pcapng



# Potential Pitfalls



- Increased CPU utilization
- Increased memory utilization
- TCP Zero Windows

